

CME 253A

# INTRODUCTION TO HIGH PERFORMANCE COMPUTING AND PARALLEL (GPU) COMPUTING

---

STANFORD SUMMER SESSION 2

26 June 2019 | Y2E2 111

Ludovic Räss

Stanford University, Department of Geophysics

[lraess@stanford.edu](mailto:lraess@stanford.edu)



# Session 2 - GPU computing

## Today's agenda

- Lecture: GPU computing and mazama GPU server
- Programming: 1/ 1D acoustic wave - Matlab vs CUDA C  
2/ Connect to mazama and run a C executable  
3/ 1D to 2D acoustic wave in CUDA C
- Tasks: 1/ 2D elastic wave  
2/ 2D viscous  
3/ towards 3D elastic or viscous

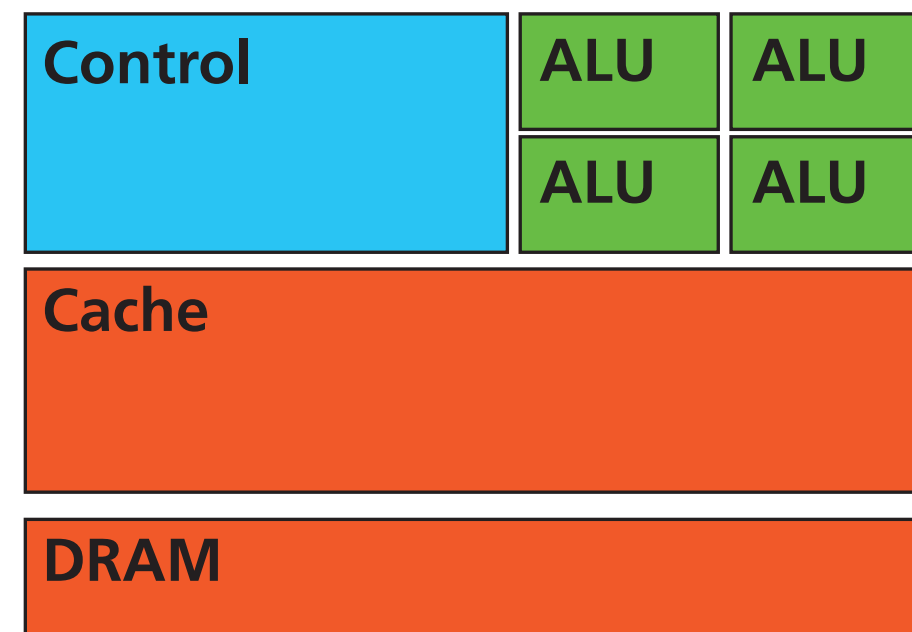
# Session 2 - GPU computing

## Today's agenda


- Lecture: GPU computing and mazama GPU server
- Programming: 1/ 1D acoustic wave - Matlab vs CUDA C  
2/ Connect to mazama and run a C executable  
3/ 1D to 2D acoustic wave in CUDA C
- Tasks: 1/ 2D elastic wave  
2/ 2D viscous  
3/ towards 3D elastic or viscous

# GPU computing


- Graphical Processing Unit (to paint the screen)
- Many small cores all doing the same job
- Large chip, close to memory
- High memory bandwidth - important for PDEs !



CPU




GPU






# Why to use a GPU

- Serial vs parallel execution > loop vs vectorised code
- Simultaneous calculations > increased concurrency
- Order of magnitude higher memory bandwidth & flop/s count



0.5 TFlop/s

50 GB/s



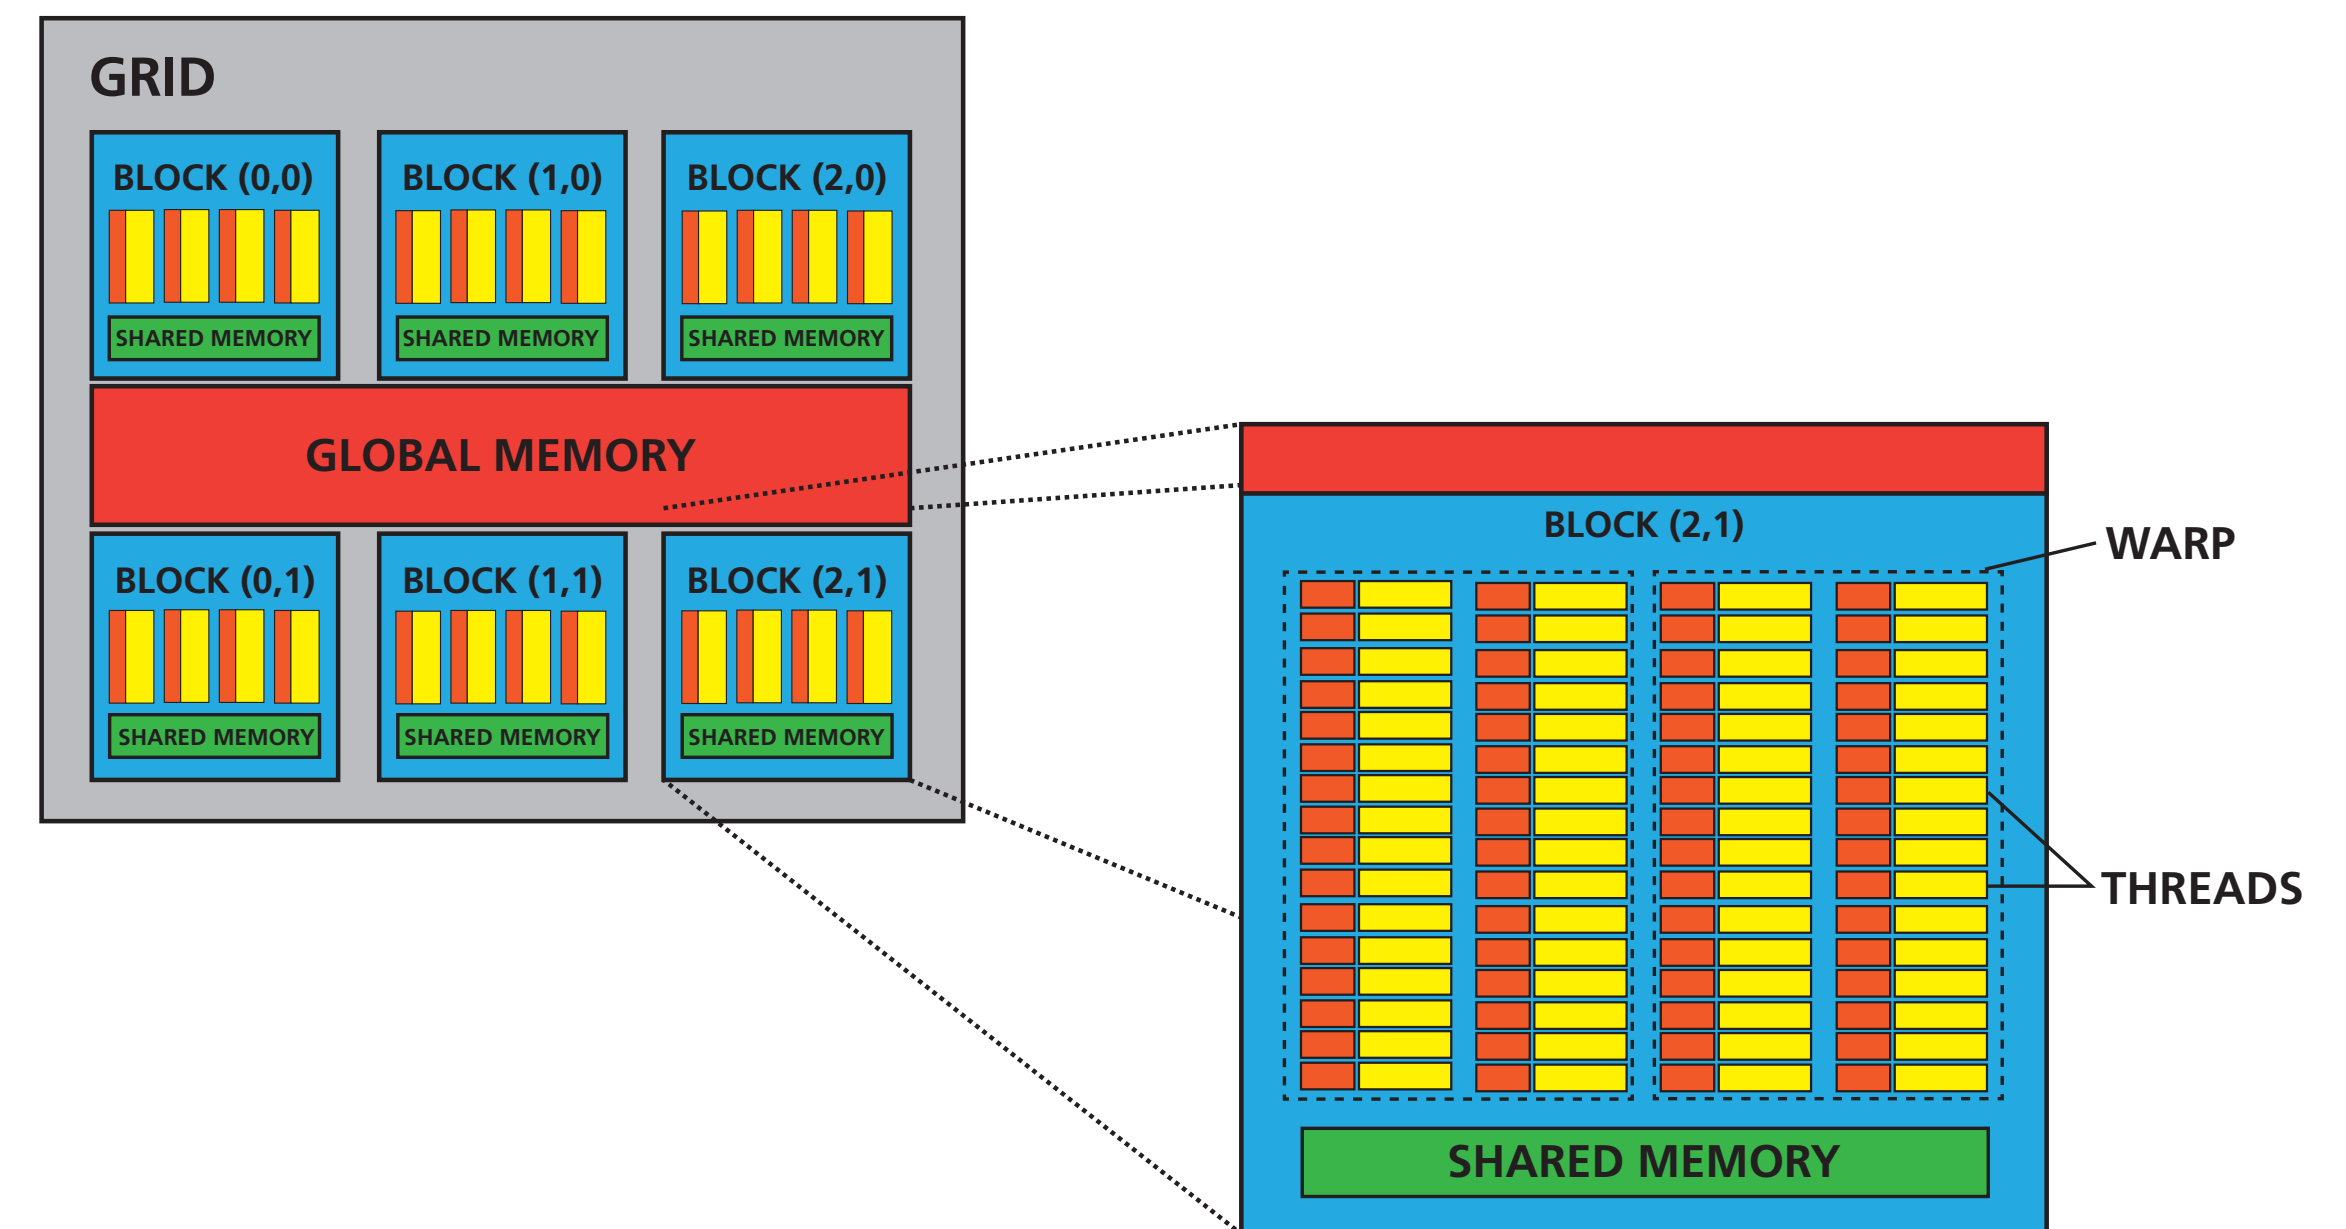
15 TFlop/s

700 GB/s

- Supercomputing available on your desk

# Programming the GPU: CUDA

- CUDA: provides vectorised indexes to C / C++ for parallel execution
- Building elements: `thread`, `block`, `grid`
- Max. 1024 threads per block !
- Assign one thread per grid point
- All threads read the same code  
loop bounds > if statements



# Session 2 - GPU computing

## Today's agenda

- Lecture: GPU computing and mazama GPU server
- Programming: 1/ 1D acoustic wave - Matlab vs CUDA C  
2/ Connect to mazama and run a C executable  
3/ 1D to 2D acoustic wave in CUDA C
- Tasks: 1/ 2D elastic wave  
2/ 2D viscous  
3/ towards 3D elastic or viscous

## 1/ 1D acoustic wave - Matlab vs CUDA C


## Acoustic wave

$$\frac{1}{k} \frac{\partial P}{\partial t} = -\frac{\partial v_x}{\partial x}$$

$$\rho \frac{\partial v_x}{\partial t} = -\frac{\partial P}{\partial x}$$

```
% Physics
Lx = 10;
k = 1;
rho = 1;
% Numerics
nx = 100;
dx = Lx/nx;
nt = 200;
dt = dx/sqrt(k/rho)/2.1;
x = (-Lx+dx)/2:dx:(Lx-dx)/2;
% Initial conditions
P = exp(-x.^2);
V = [0*P 0];
```

Initial gaussian perturbation





# 1/ 1D acoustic wave - Matlab vs CUDA C

- Minimal required changes:
  - 1/ Define some rules (macro) `#define`
  - 2/ Define `GPU_ID`
  - 3/ Define the domain size based on `grid = thread*block`
  - 4/ Copy memory from host (RAM) to device global memory
  - 5/ Define the function to run on the GPU `__global__ void`
  - 5/ Pass `block` and `grid` as argument to the GPU function
  - 7/ Gather the memory back from the device to host and visualise
- Compile using the Nvidia compiler `nvcc`

# 1/ 1D acoustic wave - Matlab vs CUDA C

- Define host (`_h`) and device (`_d`) pointers
- Initialise host: `A##_h = (DAT*)malloc( (N)*sizeof(DAT) );`  
`for(i=0; i<N; i++){A##_h[i]=(DAT)0.0;`
- Allocate device pointer: `cudaMalloc(&A##_d,N*sizeof(DAT));`
- Copy host to device:  
`cudaMemcpy(A##_d,A##_h,N*sizeof(DAT),cudaMemcpyHostToDevice);`
- Gather results:  
`cudaMemcpy(A##_h,A##_d,N*sizeof(DAT),cudaMemcpyDeviceToHost);`

# 1/ 1D acoustic wave - Matlab vs CUDA C

- Initialise vectorised grid: `dim3 grid, block;`  
Where: `block.x=BLOCK_X; grid.x=GRID_X;`
- Select the GPU:  
`gpu_id = GPU_ID; cudaSetDevice(gpu_id);`  
`cudaGetDevice(&gpu_id); cudaDeviceReset();`  
`cudaDeviceSetCacheConfig(cudaFuncCachePreferL1);`
- Reserve threads (workers) for ghost nodes: `OVERLENGTH`
- Define resolution based on grid, block: `nx = BLOCK_X*GRID_X-OVERLENGTH;`

# 1/ 1D acoustic wave - Matlab vs CUDA C

- GPU function (kernel) `__global__ void` type
- Thread index in dimension x: `int ix = blockIdx.x*blockDim.x+threadIdx.x;`
- GPU function (kernel) and device synchronisation:  
`init<<<grid,block>>>(...); cudaDeviceSynchronize();`
- Free memory host and device: `free(A##_h); cudaFree(A##_d);`

## 2/ Connect to mazama-gpu

- CEES cluster GPU resources: cees-mazama-gpu-2 and gpu-3 node
- You may need to be connected to the Stanford VPN
- Connect: `ssh <SUNetID>@cees-mazama-gpu-3`  
or `ssh <SUNetID>@cees-mazama-gpu-2` then `ssh cees-mazama-gpu-3`
- Go to scratch: `cd /scratch`
- Create a directory with your login name: `mkdir <SUNetID>`
- Copy the GPU code from your machine to mazama: `scp <source> <target>`  
`scp wave_1D.cu lraess@cees-mazama-gpu-2:/scratch/lraess`



## 2/ Connect to mazama-gpu

- `/!\ /scratch` is not backed-up. Copy your data back to your laptop or save it on data: `cd /data/cees` then `mkdir <SUNetID>`
- There is no job scheduler on mazama-gpu server nodes
  - > Always check available resources prior to run your job:
    - CPU utilisation: `top`
    - GPU utilisation: `nvidia-smi`
- Compile a GPU code using the `nvcc` compiler:  
For the Nvidia Tesla P100: `nvcc -arch=sm_70 -O3 mycode.cu`
- Execute the generated output executable: `./a.out`

# 3/ 1D to 2D acoustic wave in CUDA C

- Use linear 1D indexing in 2D and 3D:

1D:  $P(ix) \Rightarrow P[ix]$

2D:  $P(ix, iy) \Rightarrow P[ix + iy * nx]$

- Check bounds with if statements in order to exit from not active threads

Matlab loop version	CUDA C version
<pre> Vx = zeros( (nx+1)*ny, 1 ); for iyM=1:ny, iy=iyM-1;     for ix=2:nx         Vx(ix+iy*(nx+1)) = ...;     end end </pre>	<pre> zeros(Vx ,nx+1,ny ); if (iy&lt;ny &amp;&amp; ix&gt;0 &amp;&amp; ix&lt;nx){     Vx[ix + iy*(nx+1)] = ...; } </pre>

# Session 2 - GPU computing

## Today's agenda

- Lecture: GPU computing and mazama GPU server
- Programming: 1/ 1D acoustic wave - Matlab vs CUDA C  
2/ Connect to mazama and run a C executable  
3/ 1D to 2D acoustic wave in CUDA C
- Tasks: 1/ 2D elastic wave  
2/ 2D viscous (+add vertical gravity force)  
3/ towards 3D elastic or viscous

## Task 1/ 2D acoustic to elastic wave

## Acoustic waves

$$\frac{1}{k} \frac{\partial P}{\partial t} = - \left( \frac{\partial v_x}{\partial x} + \frac{\partial v_y}{\partial y} \right)$$

$$\rho \frac{\partial v_x}{\partial t} = - \frac{\partial P}{\partial x}$$

$$\rho \frac{\partial v_y}{\partial t} = - \frac{\partial P}{\partial y}$$

## Elastic waves

$$\frac{1}{k} \frac{\partial P}{\partial t} = - \left( \frac{\partial v_x}{\partial x} + \frac{\partial v_y}{\partial y} \right)$$

$$\rho \frac{\partial v_x}{\partial t} = - \frac{\partial P}{\partial x} + \frac{\partial \tau_{xx}}{\partial x} + \frac{\partial \tau_{xy}}{\partial y}$$

$$\rho \frac{\partial v_y}{\partial t} = - \frac{\partial P}{\partial y} + \frac{\partial \tau_{yy}}{\partial y} + \frac{\partial \tau_{xy}}{\partial x}$$

$$\frac{\partial \tau_{xx}}{\partial t} = 2G \left( \frac{\partial v_x}{\partial x} - \frac{1}{3} \left( \frac{\partial v_x}{\partial x} + \frac{\partial v_y}{\partial y} \right) \right)$$

$$\frac{\partial \tau_{yy}}{\partial t} = 2G \left( \frac{\partial v_y}{\partial y} - \frac{1}{3} \left( \frac{\partial v_x}{\partial x} + \frac{\partial v_y}{\partial y} \right) \right)$$

$$\frac{\partial \tau_{xy}}{\partial t} = 2G \frac{1}{2} \left( \frac{\partial v_x}{\partial y} + \frac{\partial v_y}{\partial x} \right)$$

## Task 2/ 2D elastic to viscous

Elastic waves	Viscous flow
$\frac{1}{k} \frac{\partial P}{\partial t} = - \left( \frac{\partial v_x}{\partial x} + \frac{\partial v_y}{\partial y} \right)$	$\frac{1}{k} \frac{\partial P}{\partial t} = - \left( \frac{\partial v_x}{\partial x} + \frac{\partial v_y}{\partial y} \right)$
$\rho \frac{\partial v_x}{\partial t} = - \frac{\partial P}{\partial x} + \frac{\partial \tau_{xx}}{\partial x} + \frac{\partial \tau_{xy}}{\partial y}$	$\rho \frac{\partial v_x}{\partial t} = - \frac{\partial P}{\partial x} + \frac{\partial \tau_{xx}}{\partial x} + \frac{\partial \tau_{xy}}{\partial y}$
$\rho \frac{\partial v_y}{\partial t} = - \frac{\partial P}{\partial y} + \frac{\partial \tau_{yy}}{\partial y} + \frac{\partial \tau_{xy}}{\partial x}$	$\rho \frac{\partial v_y}{\partial t} = - \frac{\partial P}{\partial y} + \frac{\partial \tau_{yy}}{\partial y} + \frac{\partial \tau_{xy}}{\partial x}$
$\frac{\partial \tau_{xx}}{\partial t} = 2G \left( \frac{\partial v_x}{\partial x} - \frac{1}{3} \left( \frac{\partial v_x}{\partial x} + \frac{\partial v_y}{\partial y} \right) \right)$	$\tau_{xx} = 2\eta \left( \frac{\partial v_x}{\partial x} - \frac{1}{3} \left( \frac{\partial v_x}{\partial x} + \frac{\partial v_y}{\partial y} \right) \right)$
$\frac{\partial \tau_{yy}}{\partial t} = 2G \left( \frac{\partial v_y}{\partial y} - \frac{1}{3} \left( \frac{\partial v_x}{\partial x} + \frac{\partial v_y}{\partial y} \right) \right)$	$\tau_{yy} = 2\eta \left( \frac{\partial v_y}{\partial y} - \frac{1}{3} \left( \frac{\partial v_x}{\partial x} + \frac{\partial v_y}{\partial y} \right) \right)$
$\frac{\partial \tau_{xy}}{\partial t} = 2G \frac{1}{2} \left( \frac{\partial v_x}{\partial y} + \frac{\partial v_y}{\partial x} \right)$	$\tau_{xy} = 2\eta \frac{1}{2} \left( \frac{\partial v_x}{\partial y} + \frac{\partial v_y}{\partial x} \right)$



# Outlook - Session 3

- Topic: Performance evaluation
- Programming: Time measurement, GB/s vs Flops/s, effective memory throughput
- Tasks: 2D and 3D elastic waves or viscous Stokes
- Mandatory for session 3: Matlab code in both loop and vectorised style. CUDA C GPU code reproducing Matlab results.

That's it for today

- ● ● [wp.unil.ch/geocomputing/](http://wp.unil.ch/geocomputing/)

- ● ● [Iraess@stanford.edu](mailto:Iraess@stanford.edu)

