

CME 253A

# INTRODUCTION TO HIGH PERFORMANCE COMPUTING AND PARALLEL (GPU) COMPUTING

---

STANFORD SUMMER SESSION 5

15 July 2019 | Y2E2 111

Ludovic Räss

Stanford University, Department of Geophysics

[lraess@stanford.edu](mailto:lraess@stanford.edu)



# Session 5 - CUDA MPI and Q&A

## Today's agenda

- Lecture: Multi-GPU implementation with CUDA MPI
- Programming: 1/ Conceptual MATLAB example  
2/ CUDA MPI 1D wave propagation
- Q&A regarding the final project

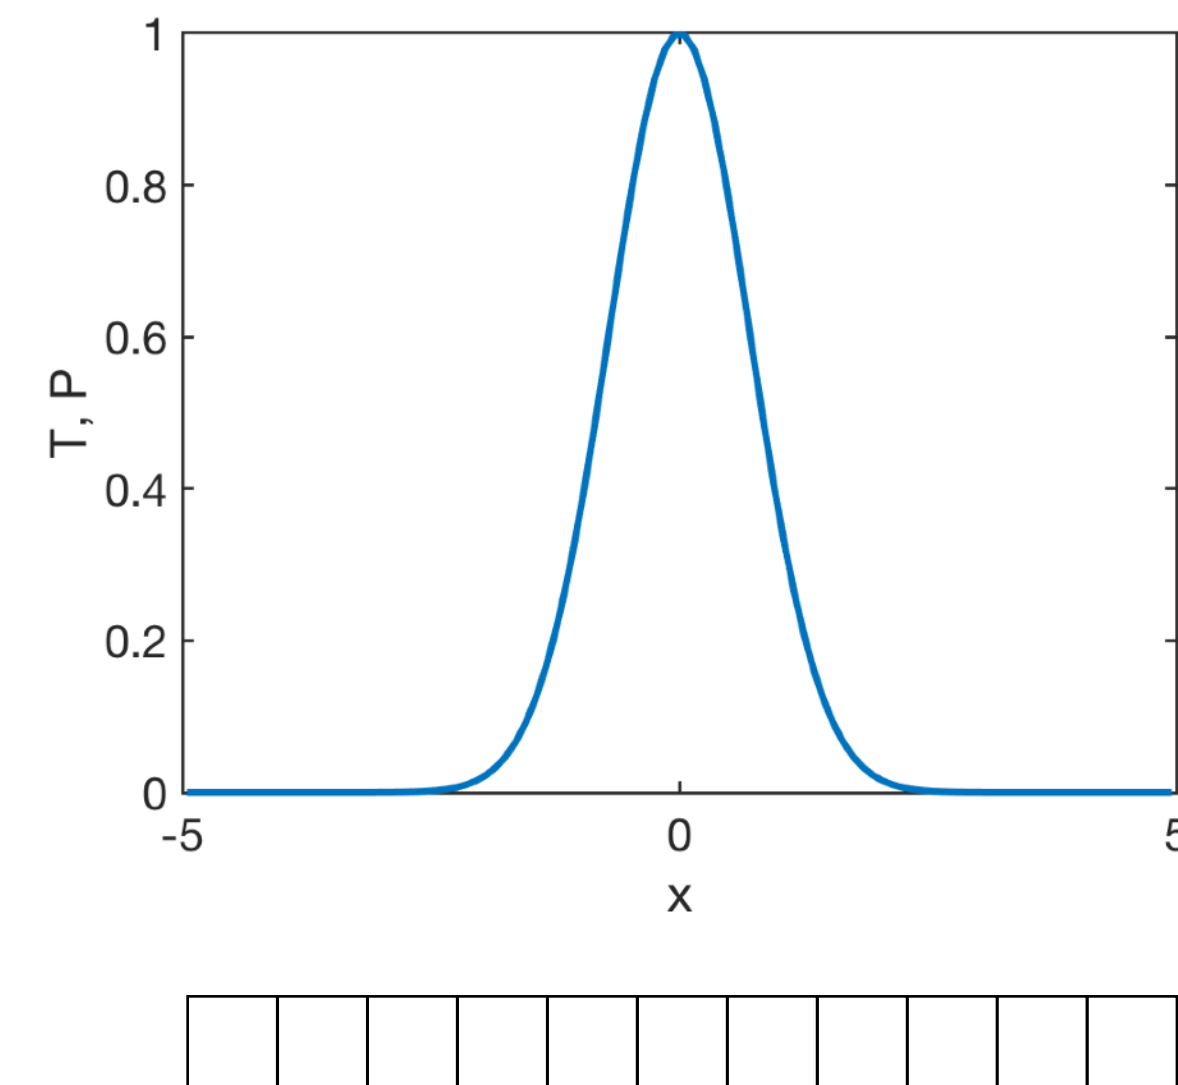
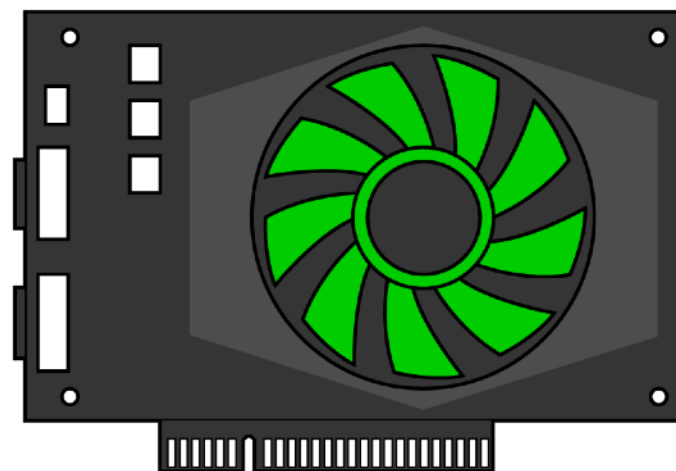
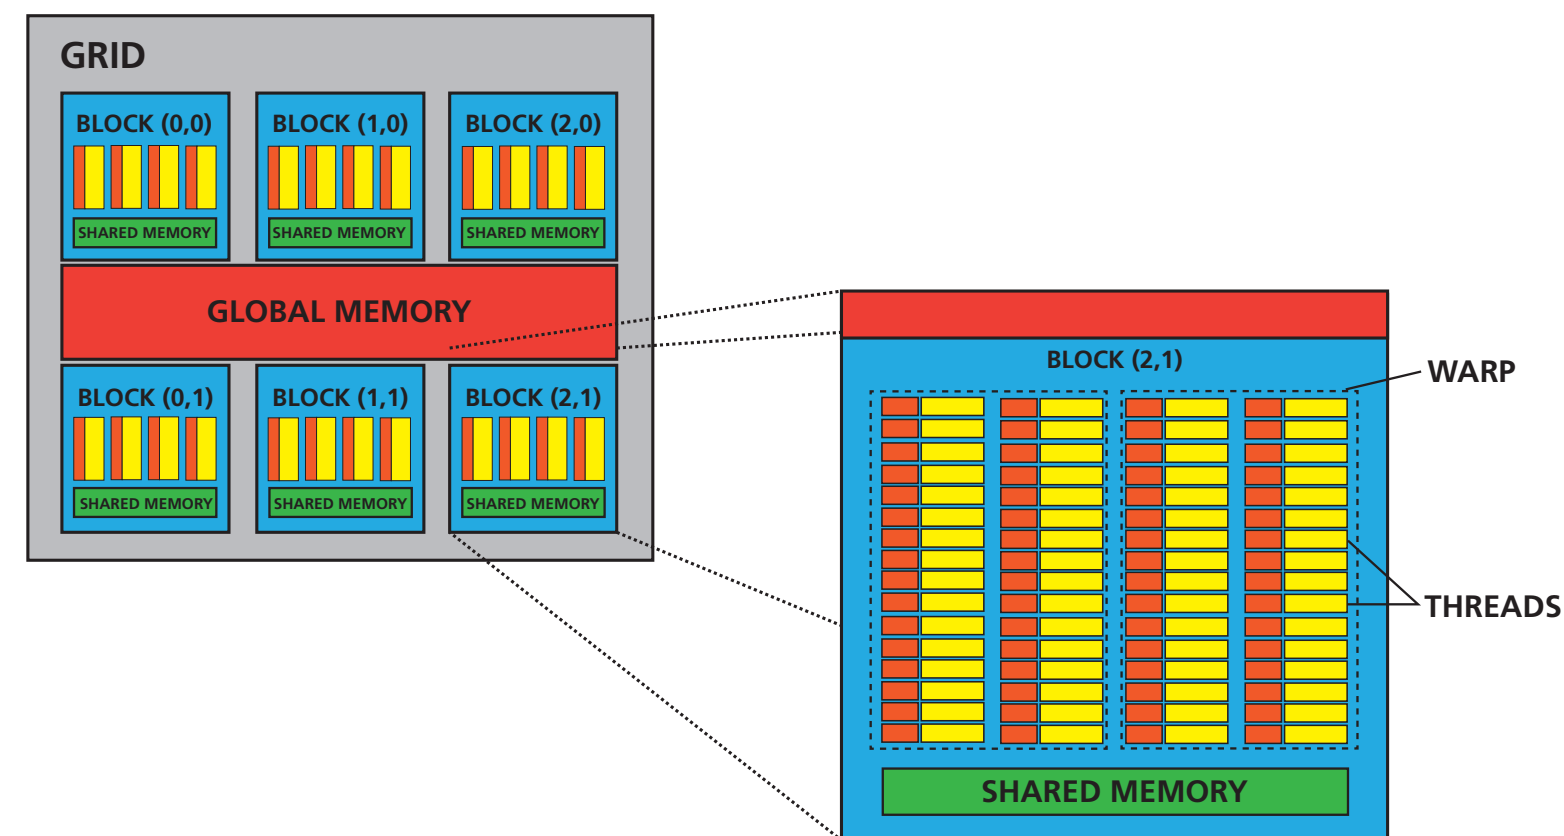
# Session 5 - CUDA MPI and Q&A

## Today's agenda

- Lecture: Multi-GPU implementation with CUDA MPI
- Programming: 1/ Conceptual MATLAB example  
2/ CUDA MPI 1D wave propagation
- Q&A regarding the final project

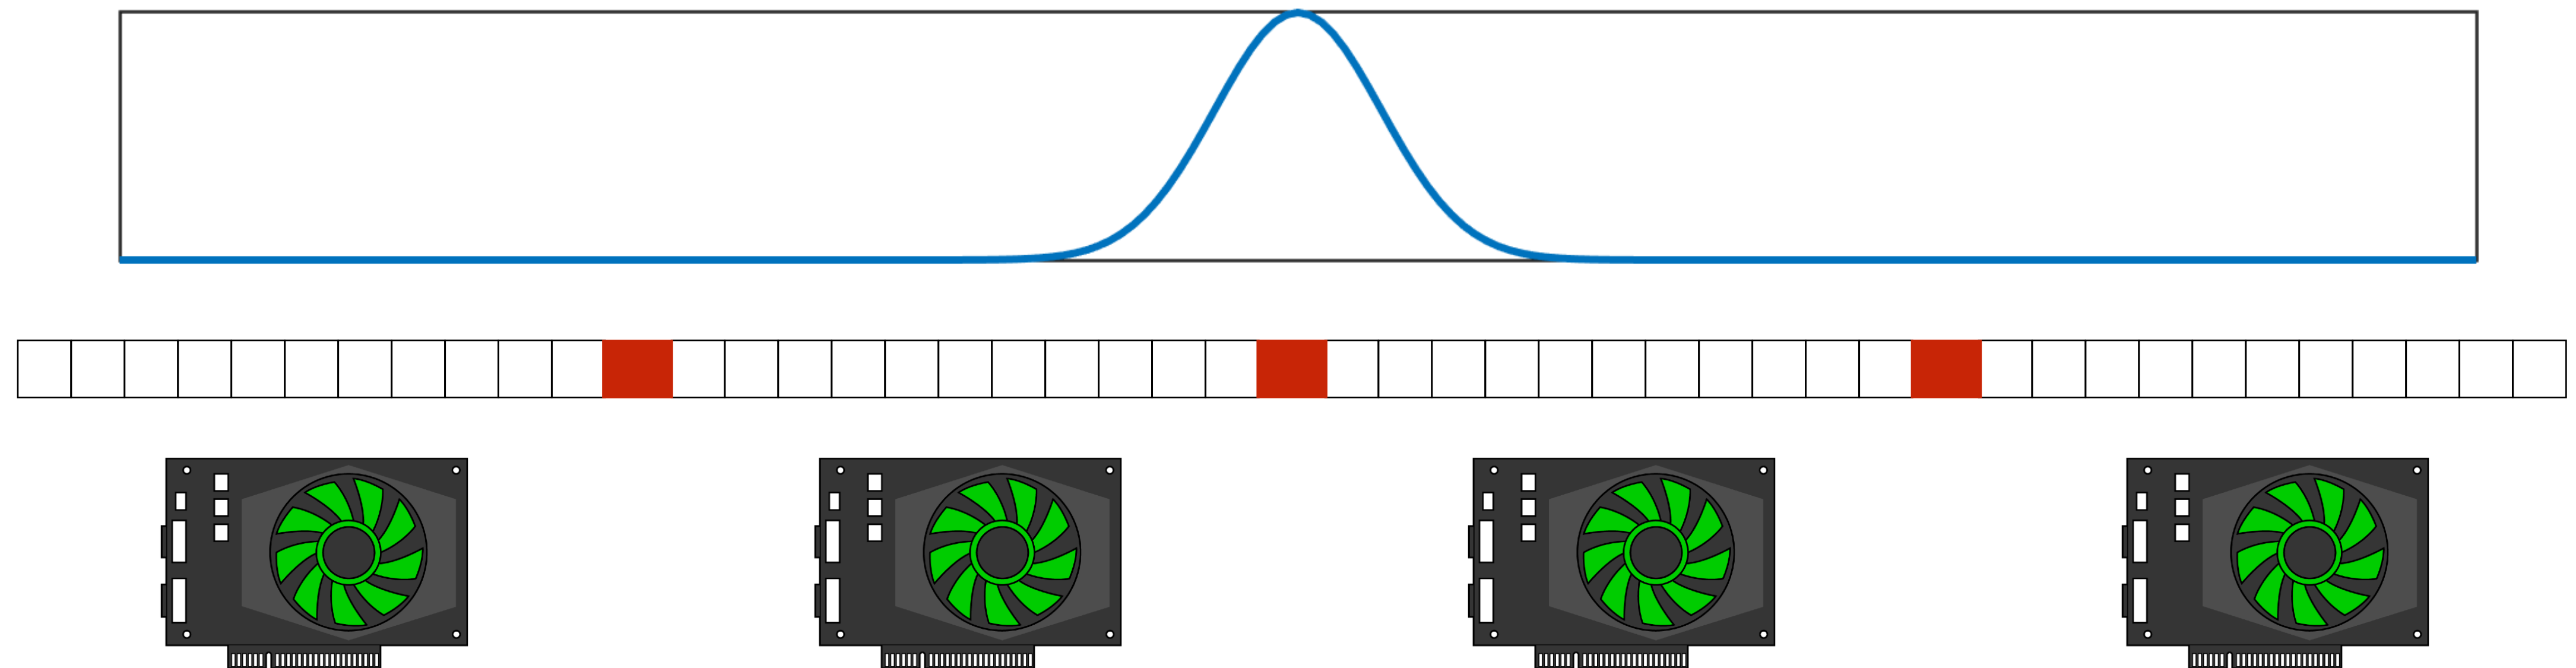
# Multi-GPU implementation with CUDA MPI

- Up to now we mapped and parallelised our FD domain on a single GPU
- All threads share the same global GPU memory



# Multi-GPU implementation with CUDA MPI

- How to use 2+ GPUs in order to compute larger domains
- Using multiple GPUs requires to handle distributed memory
- MPI (message passing interface) library is useful in this task



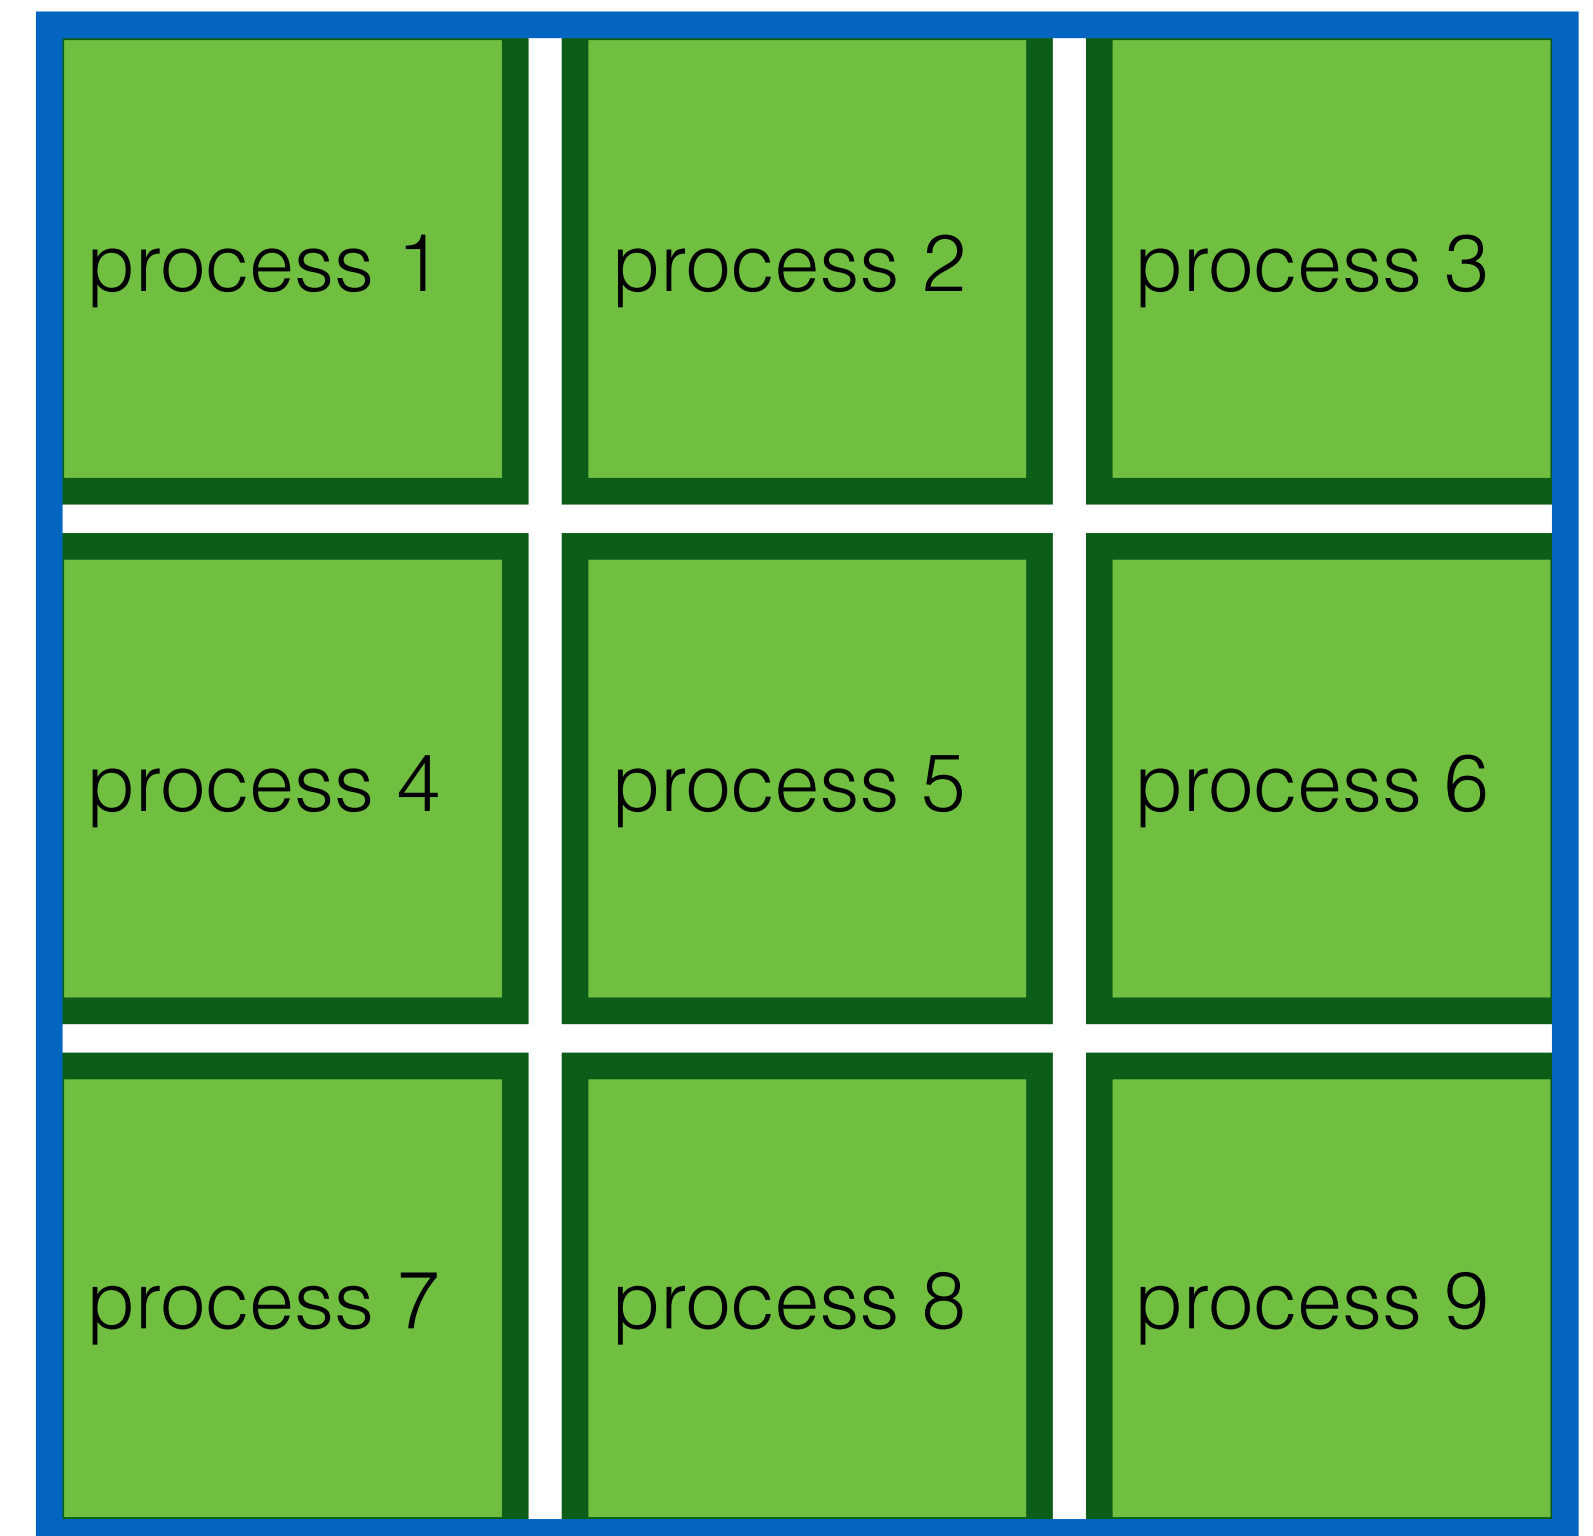
# Multi-GPU implementation with CUDA MPI

- Global domain
- Local processes
- Boundary conditions:
  - Global > physics
  - Local > MPI

Conceptual 2-D domain decomposition

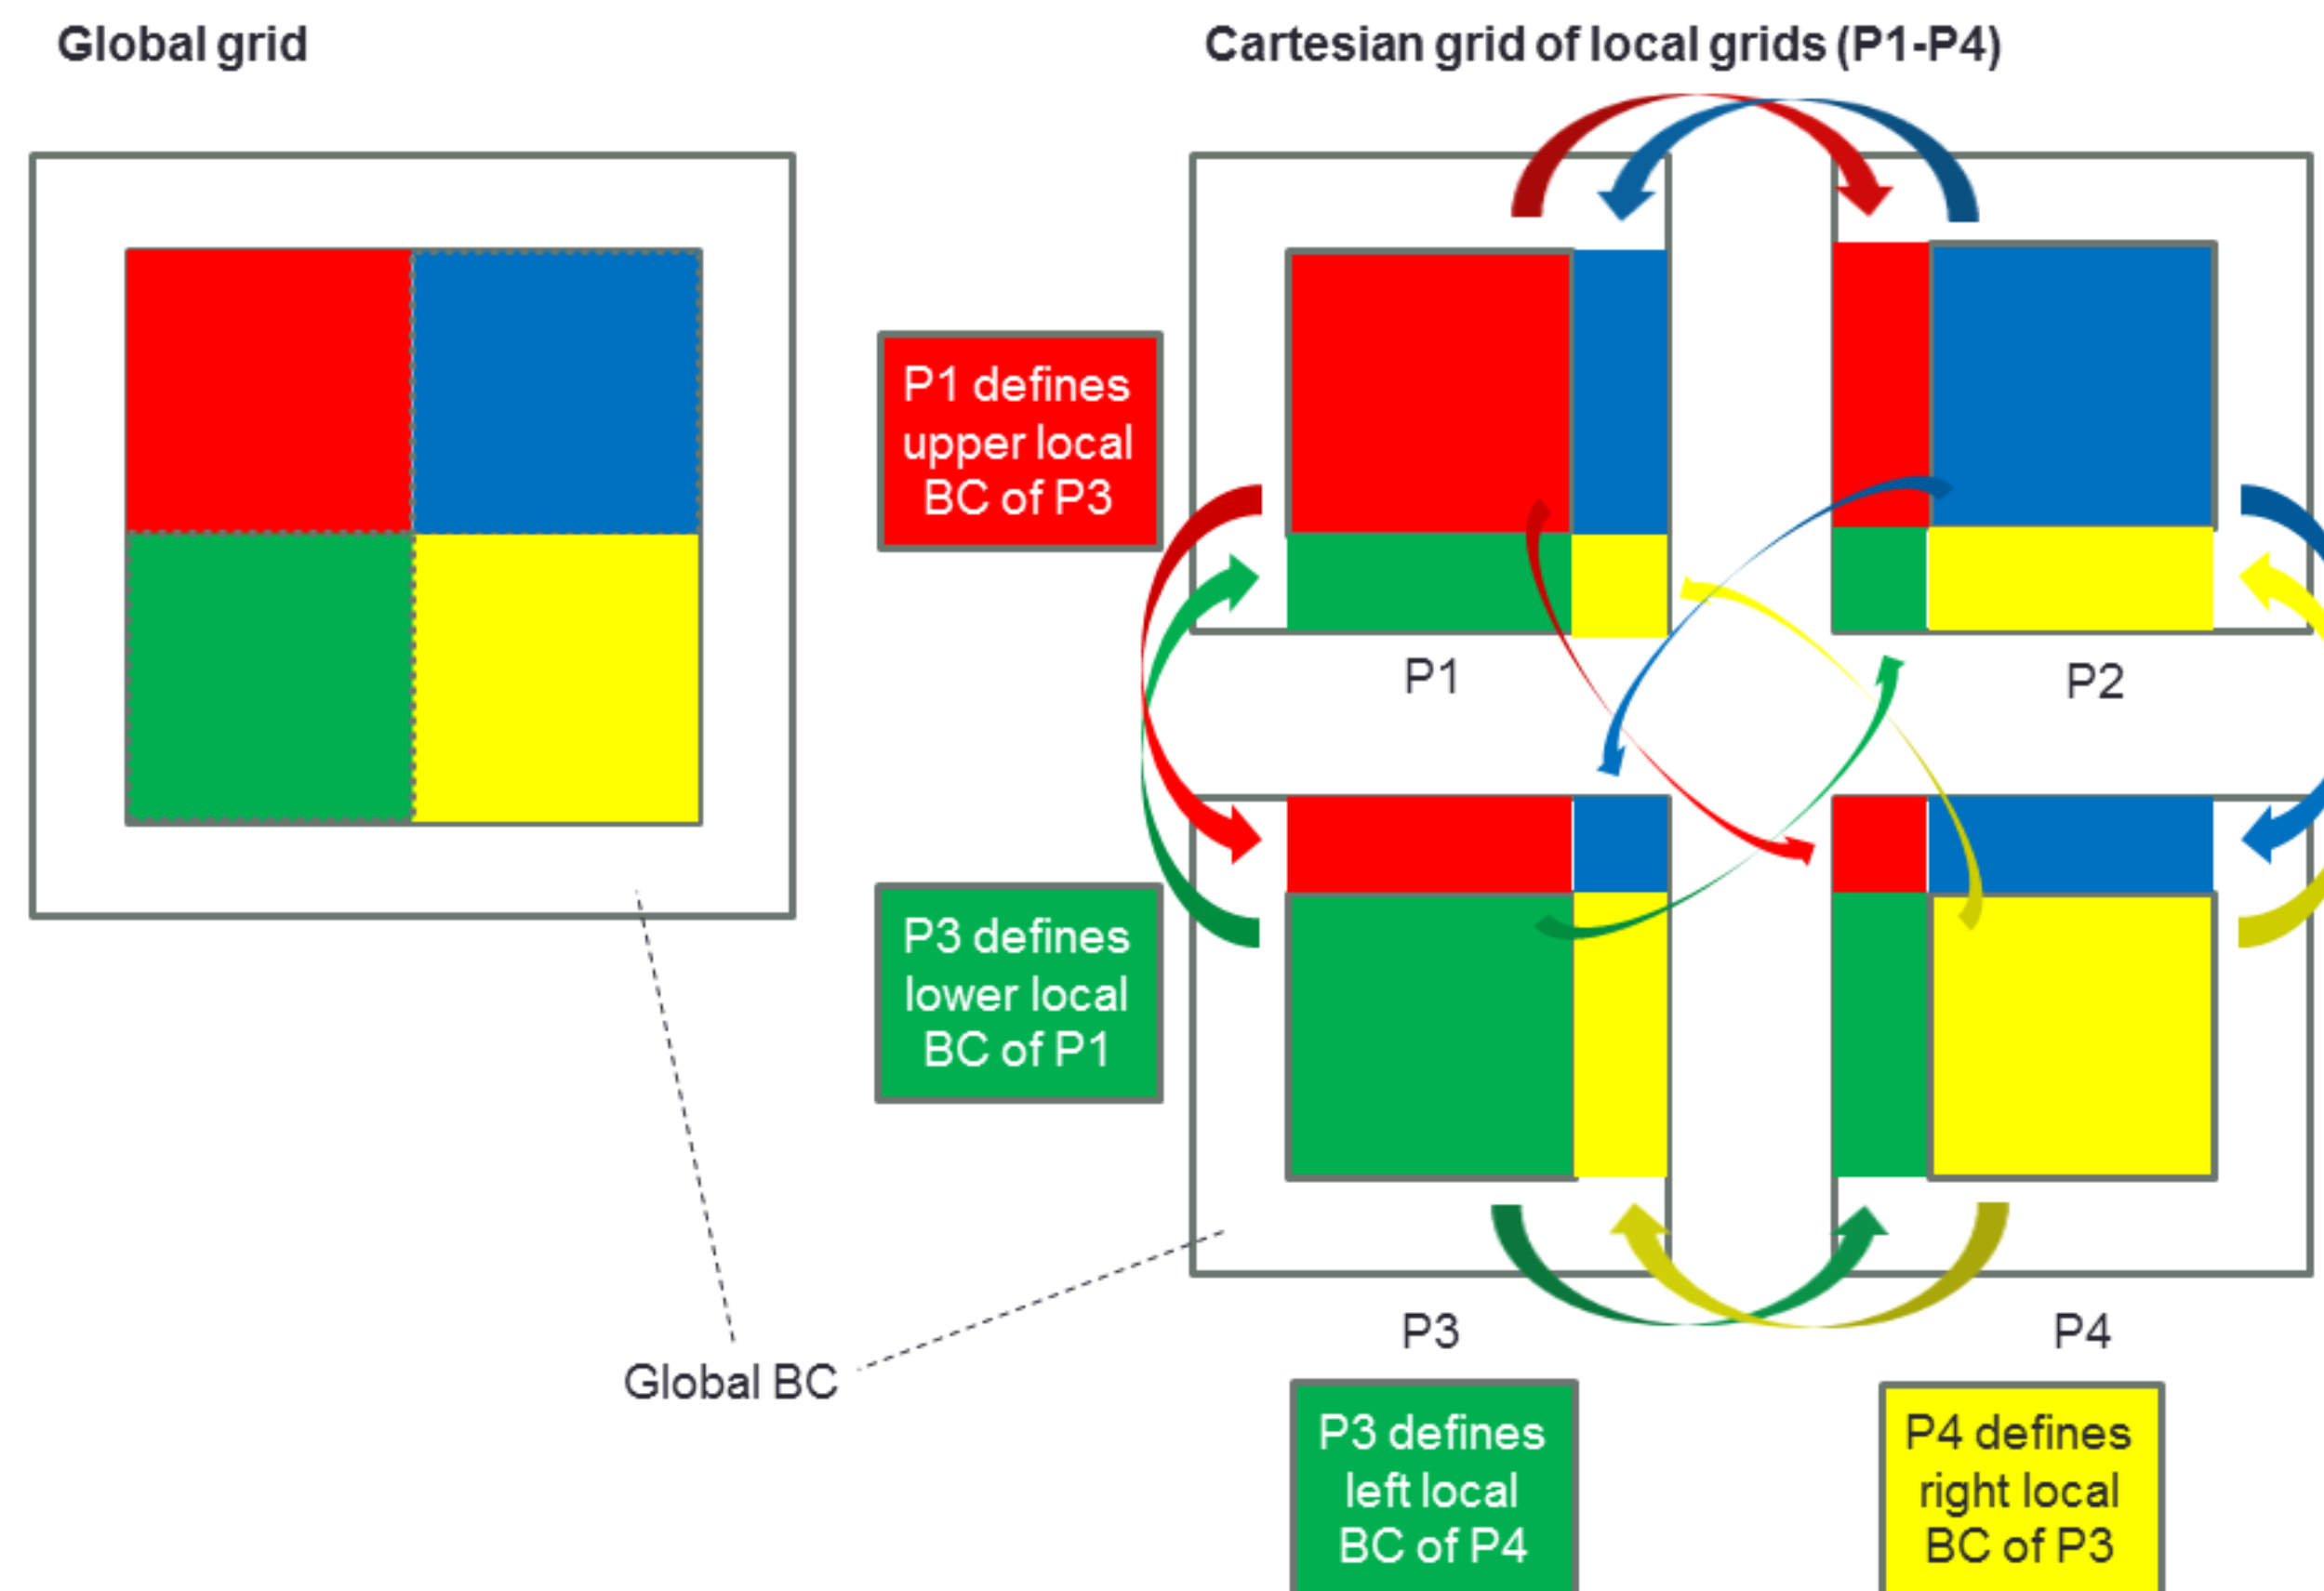


Conceptual 2-D domain decomposition



# Multi-GPU implementation with CUDA MPI

- All about exchanging the boundary values



# Multi-GPU implementation with CUDA MPI

MPI - a CPU library that has (a priori) nothing to do with GPU knowledge

- Use MPI to assign one GPU to each MPI process
- “Vectorised” domain hierarchy: **CUDA Threads, Blocks, Grid + MPI Dims**
- Naive design: - use MPI rank to get gpu\_id (get local rank per node for correct gpu\_id)
  - exchange inner boundary conditions between neighbouring processes
  - send and receive host pointers via MPI, needs both CPU and GPU buffers



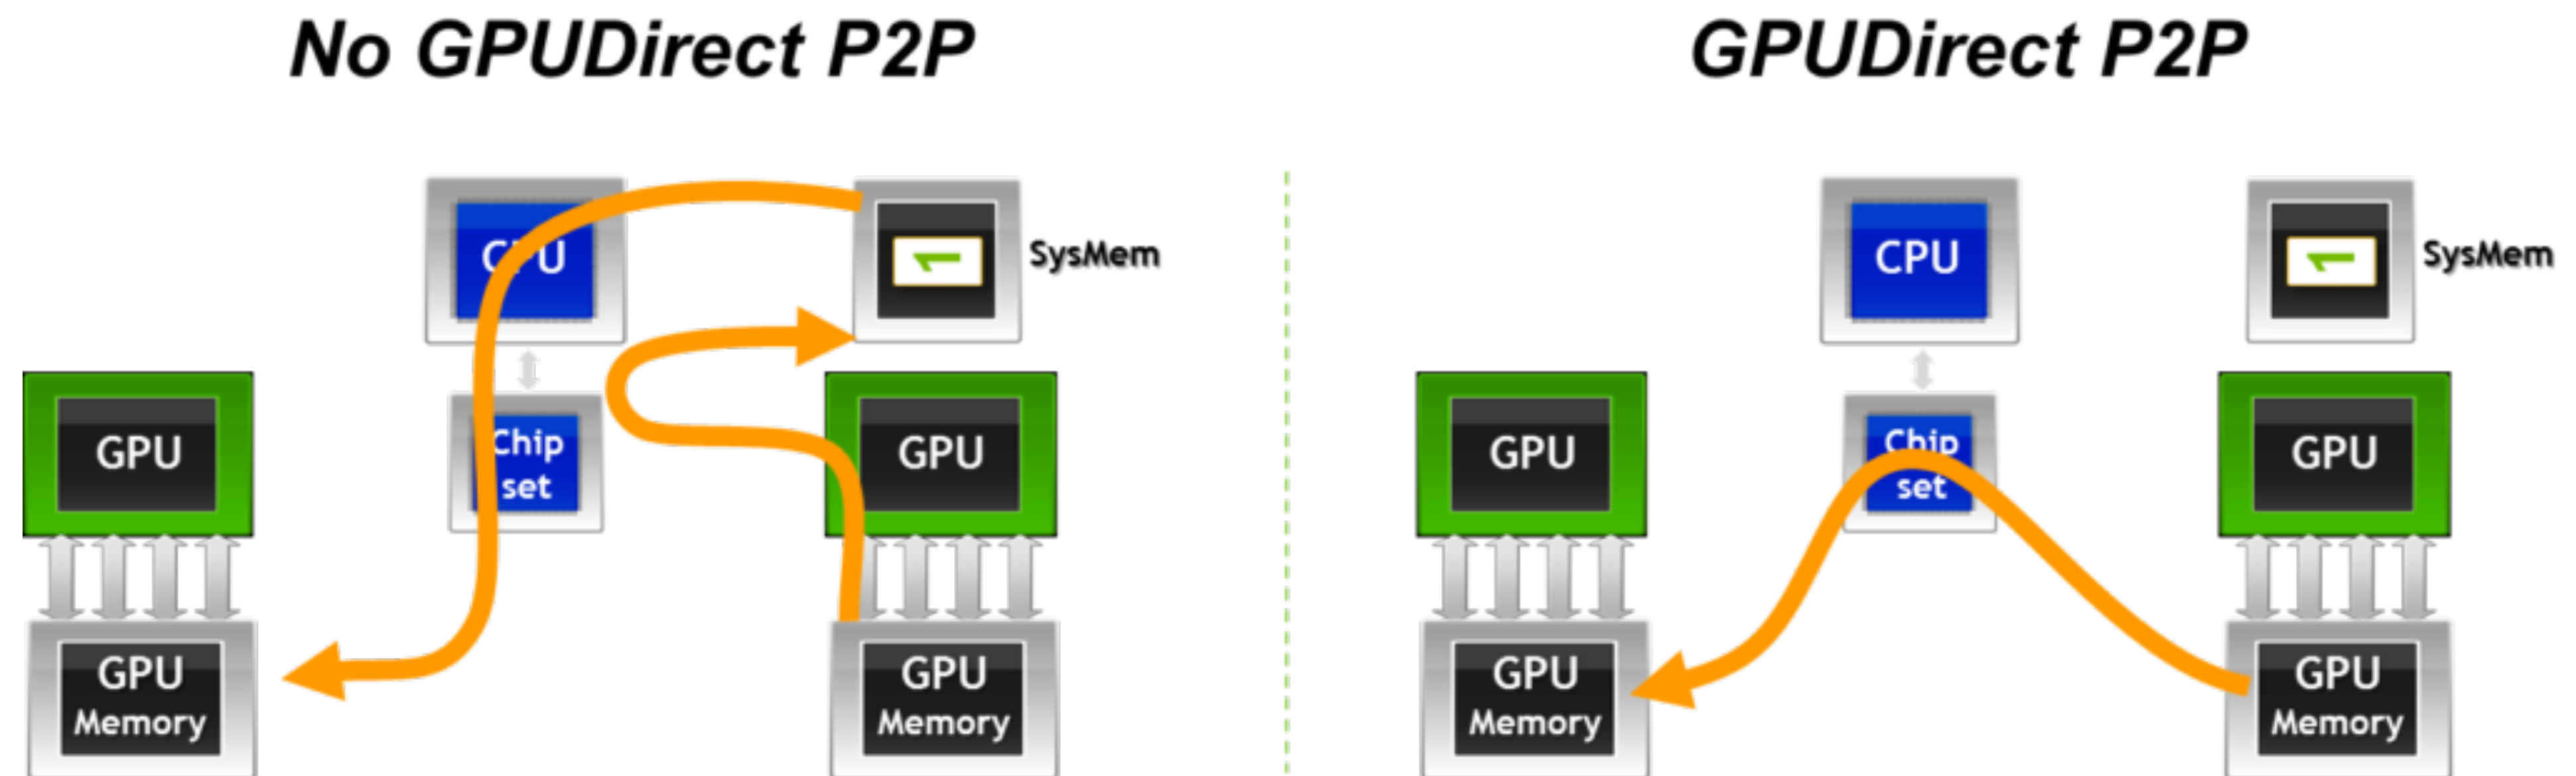
# Multi-GPU implementation with CUDA MPI

## CUDA-aware MPI

- MPI send and receive functions can accept device pointers (GPU)
- No need for explicit buffer copy on the host
- CPU and GPU MPI are nearly identical except:  
Context initialisation  
Pointers (host or device) passed to the send/recv MPI functions

# Multi-GPU implementation with CUDA MPI

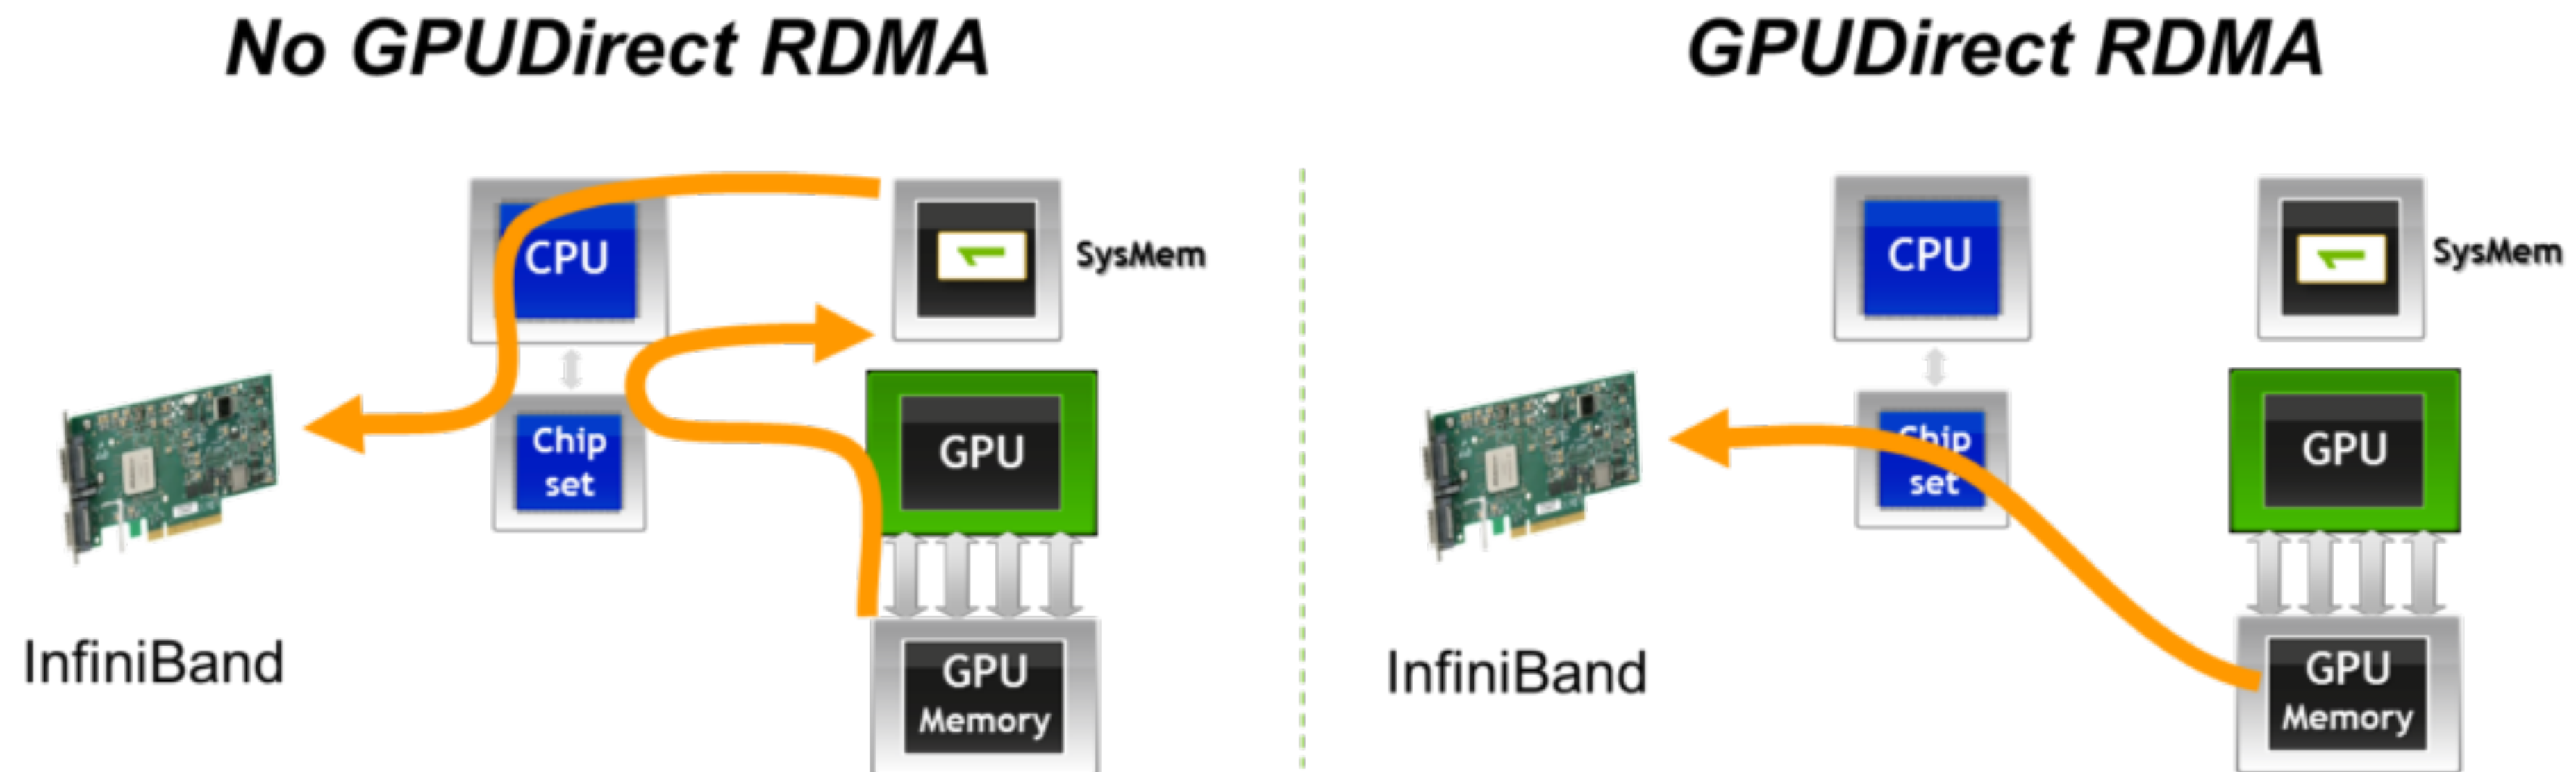
- GPU + MPI - what is going on within a multi-GPU node (Peer-to-peer P2P)



<https://devblogs.nvidia.com/introduction-cuda-aware-mpi/>

# Multi-GPU implementation with CUDA MPI

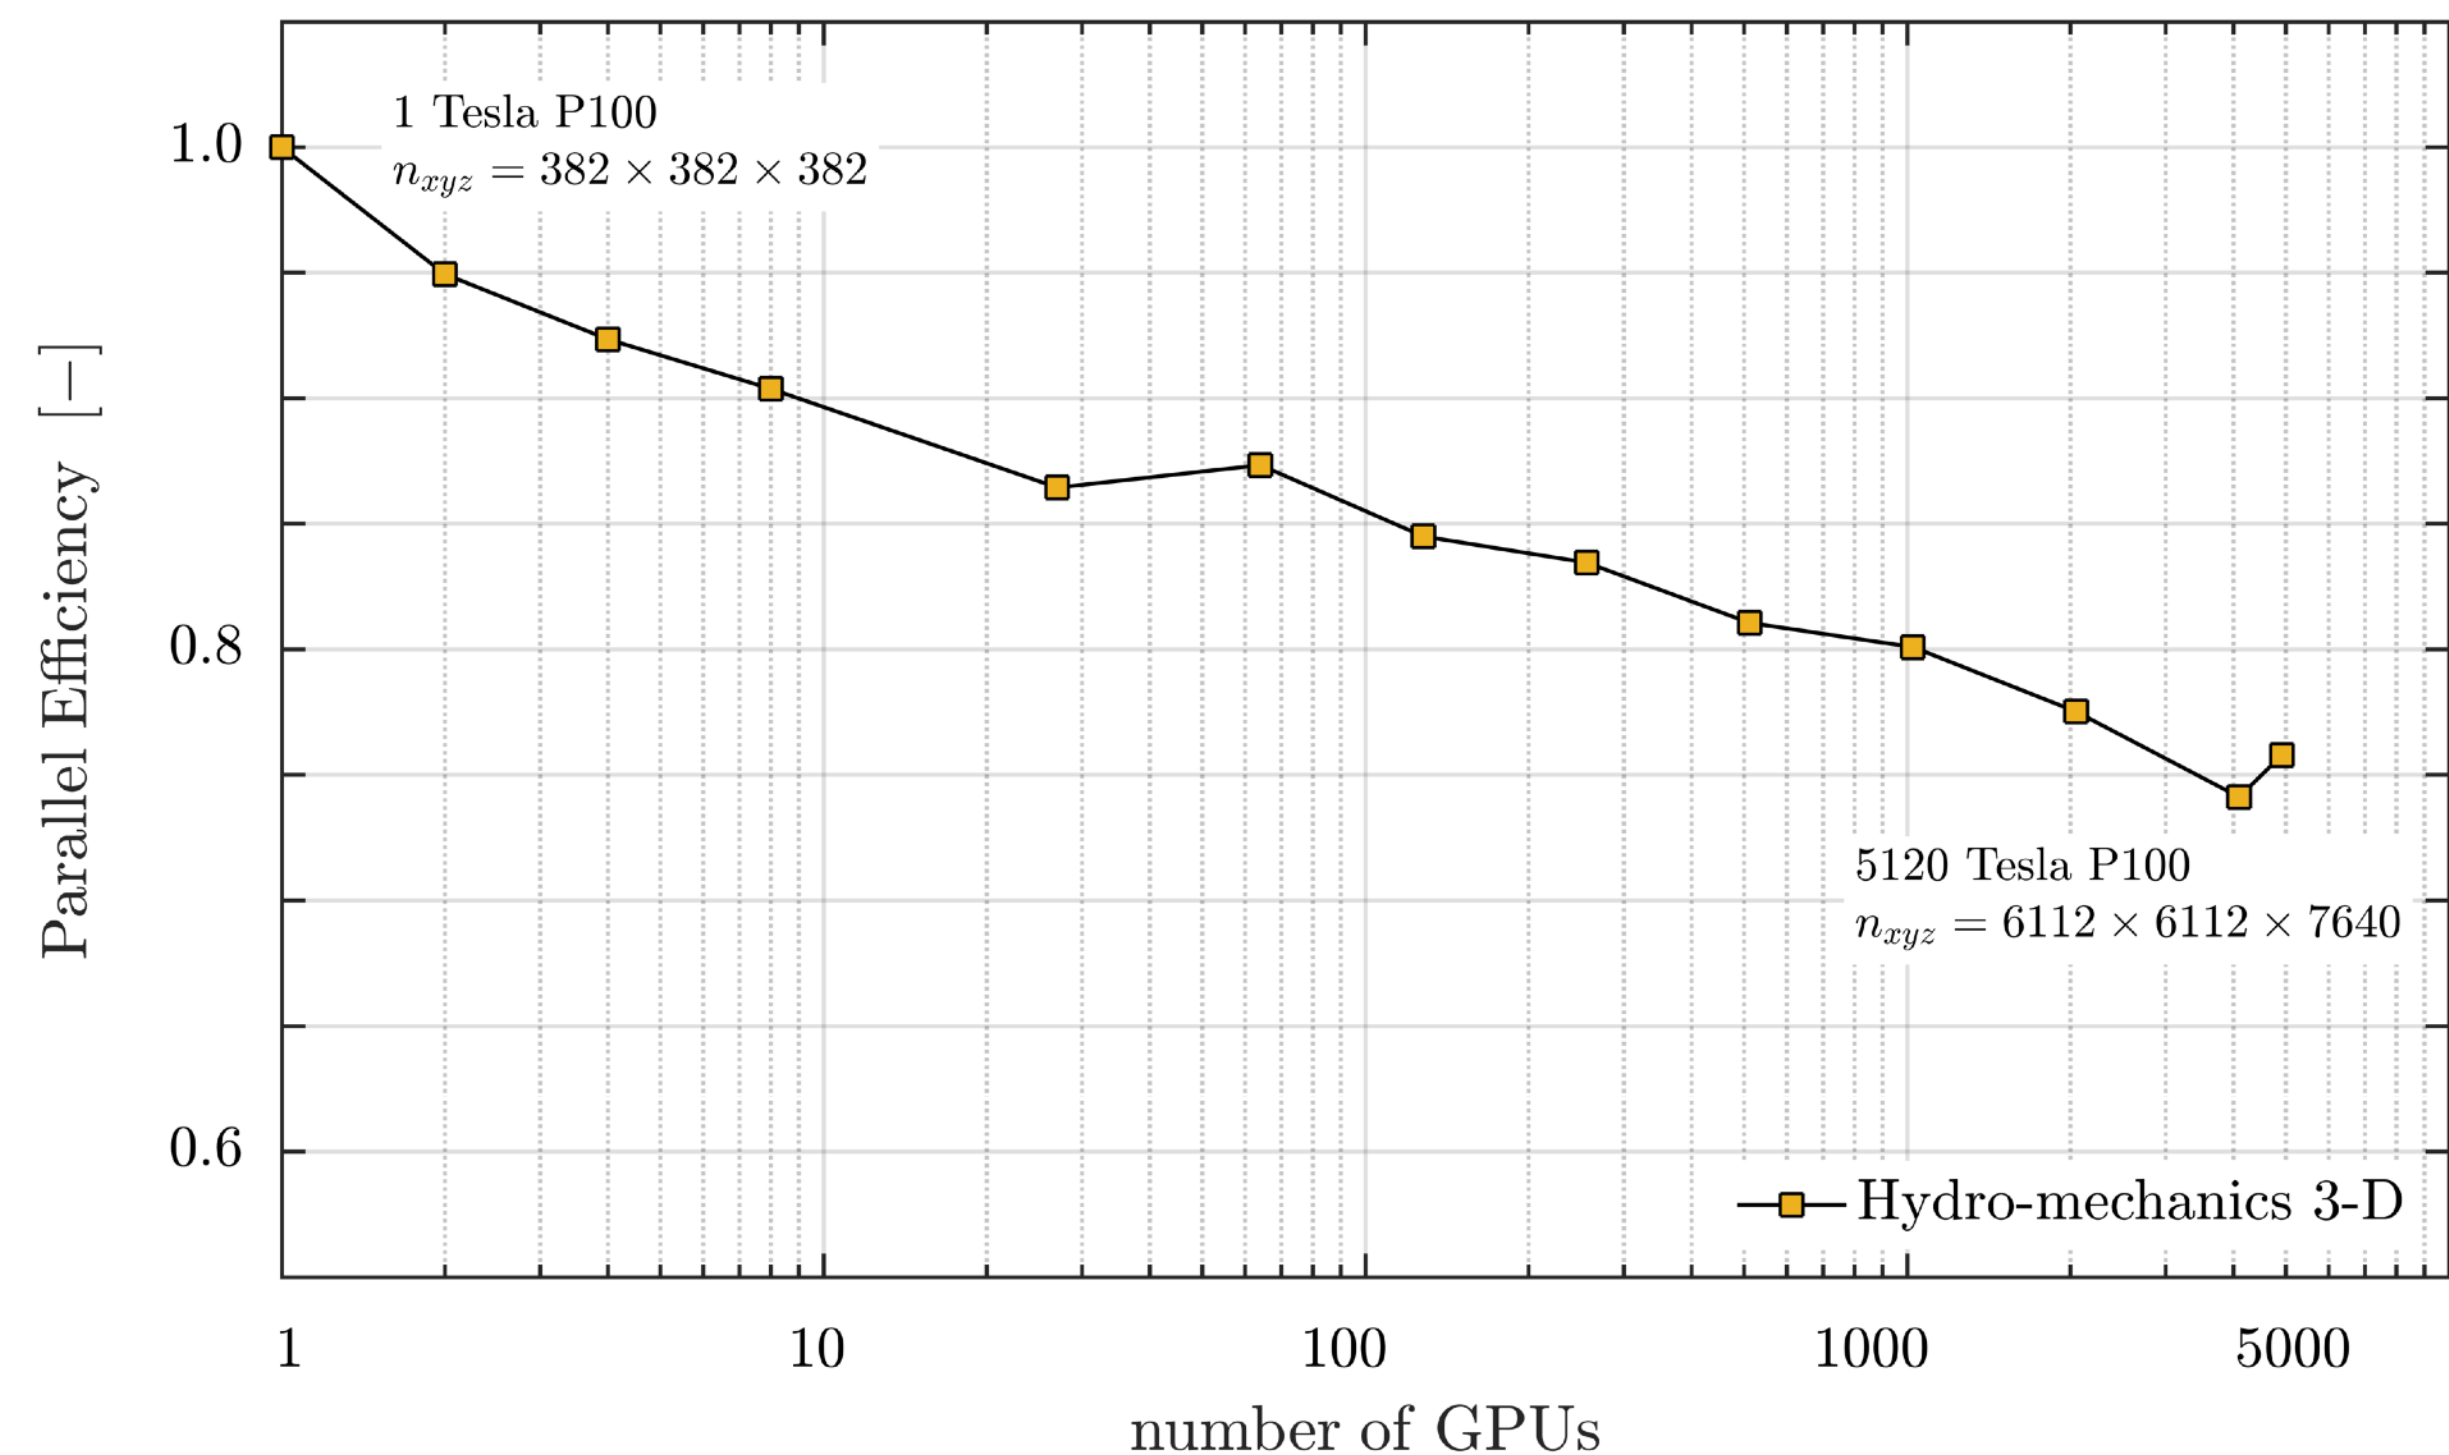
- GPU + MPI - what is going on between different GPU nodes (RDMA + IB)



<https://devblogs.nvidia.com/introduction-cuda-aware-mpi/>

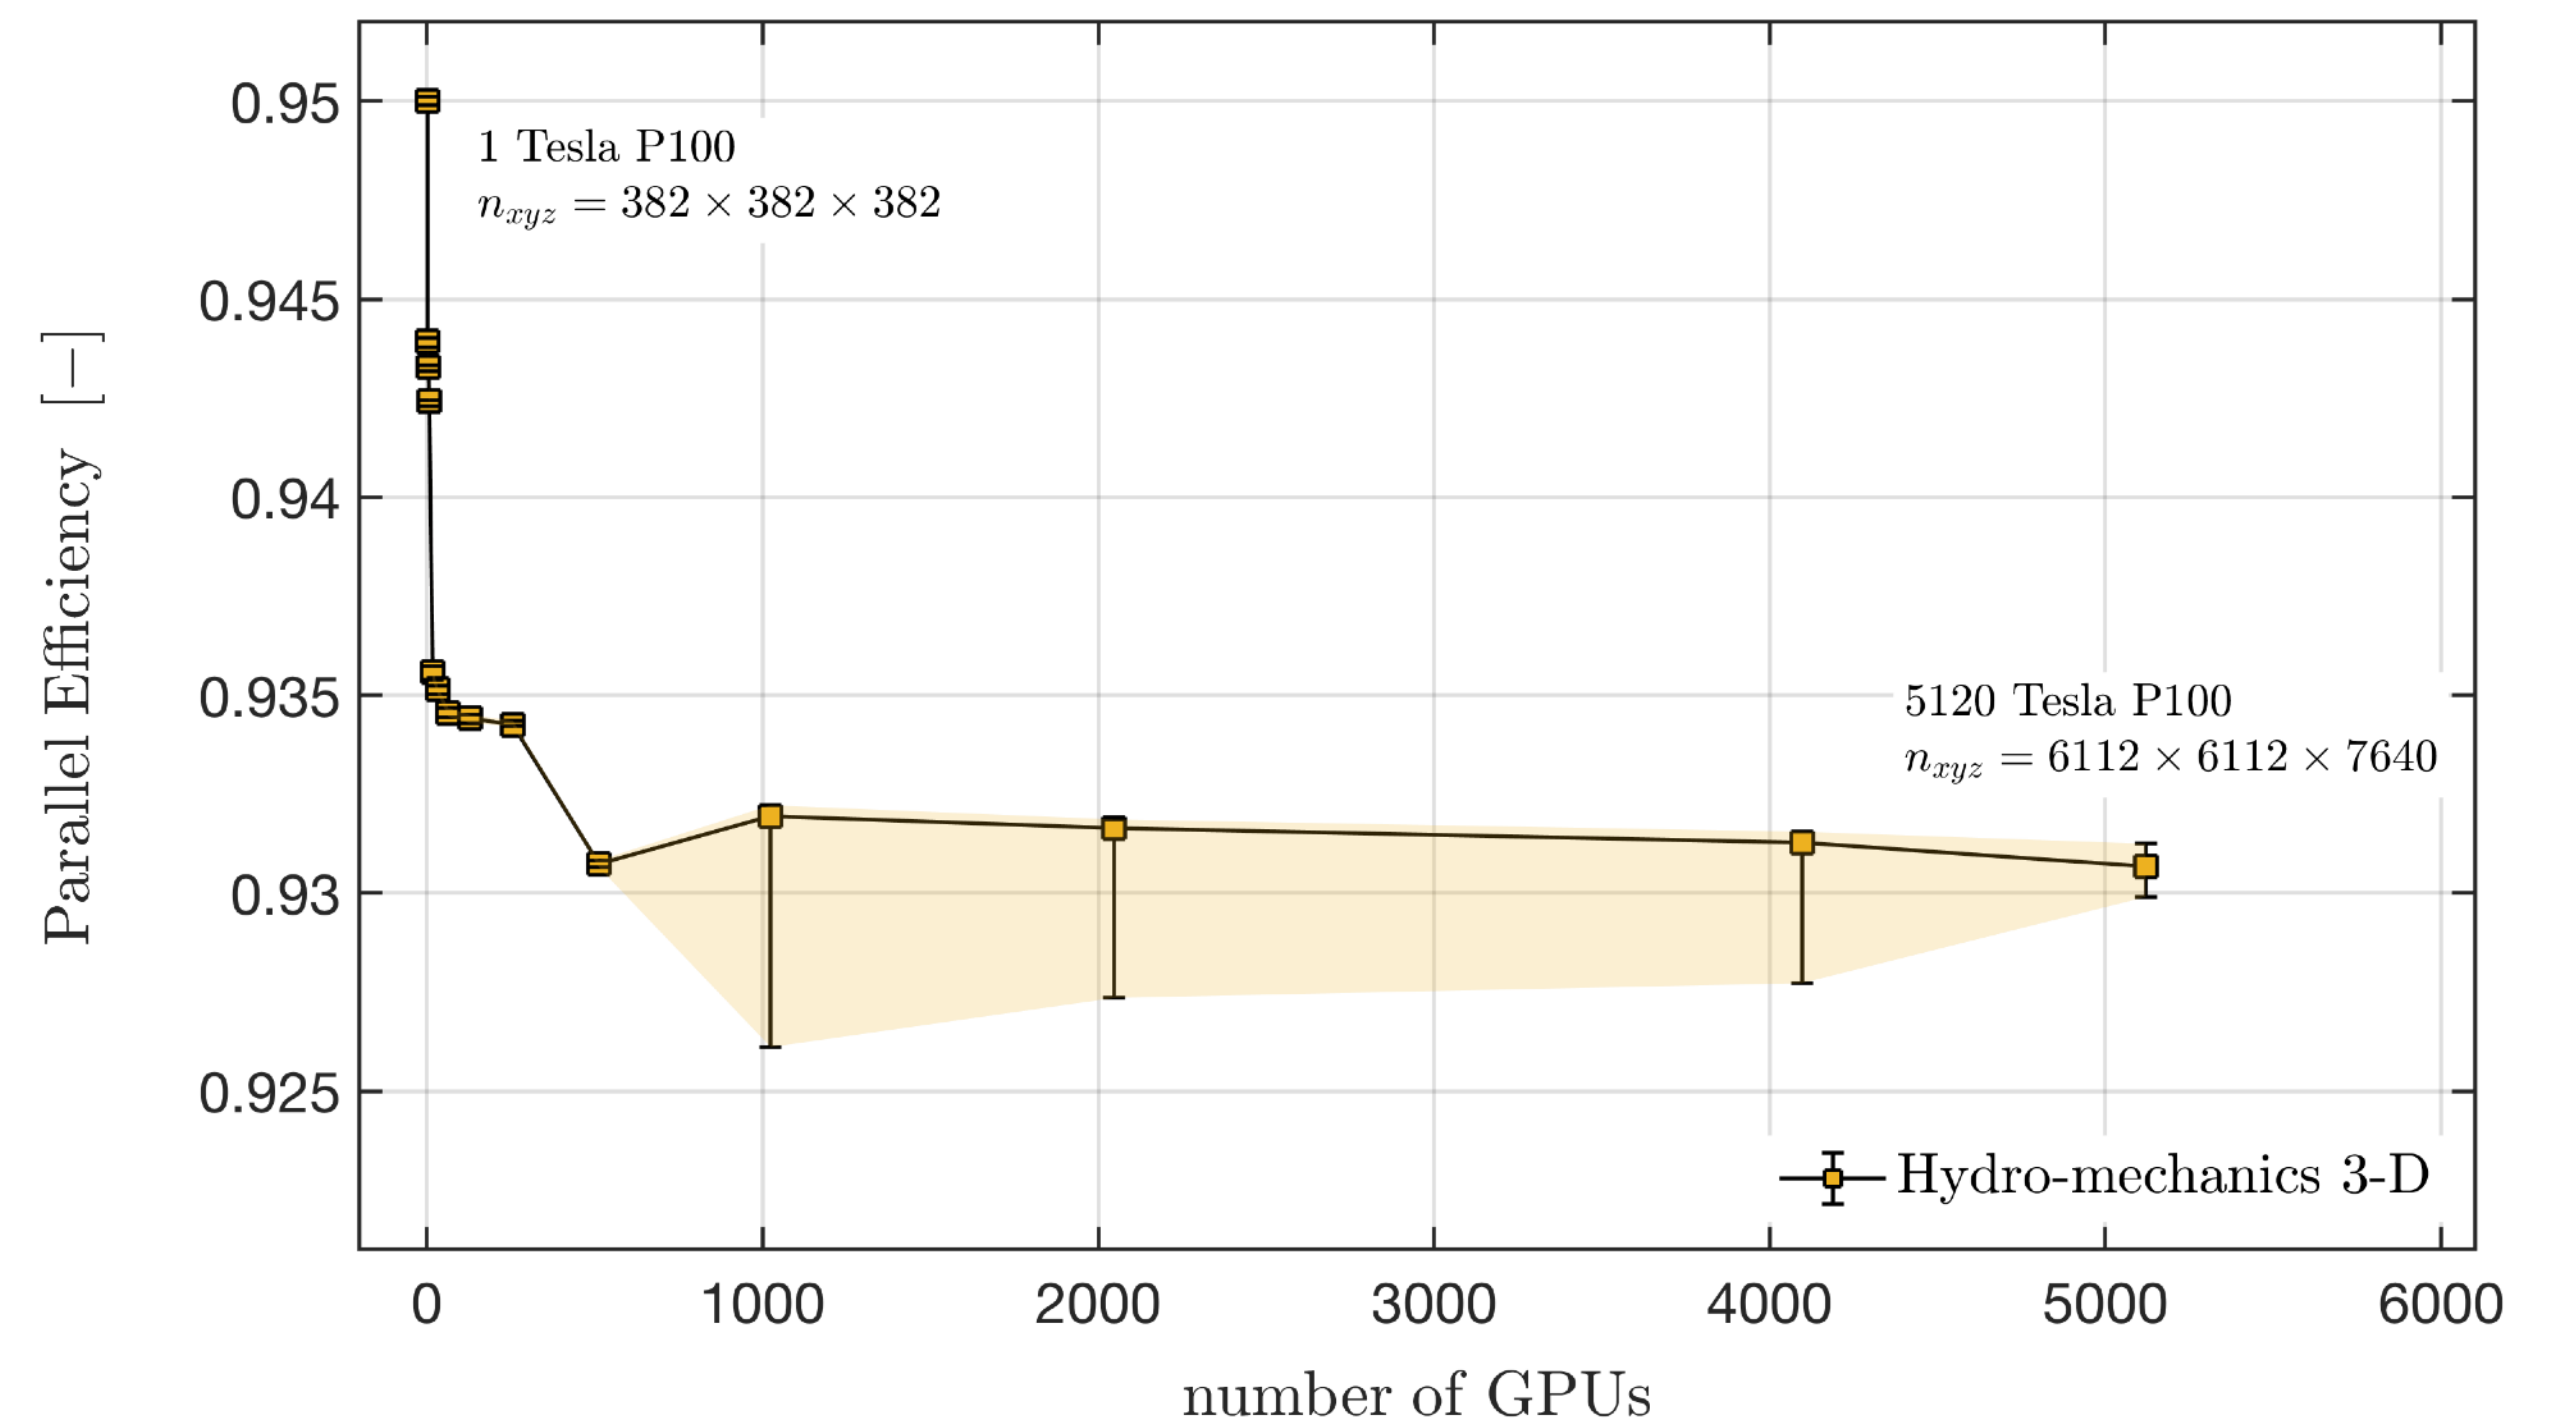
# Multi-GPU implementation with CUDA MPI

- Parallel efficiency
- How much time is lost in communication
- Ideal: close to 1
- MPI point-to-point shows ok perfs
- But perfs gets worse with growing number of MPI processes
- Good news - there is a fix !



# Multi-GPU implementation with CUDA MPI

- Parallel efficiency close to 1
- Hiding MPI communication with computations
- Straight forward implementation in using CUDA streams
- Works well with stencil codes and domain decomposition





# Session 5 - CUDA MPI and Q&A

## Today's agenda

- Lecture: Multi-GPU implementation with CUDA MPI
- Programming: 1/ Conceptual MATLAB example  
2/ CUDA MPI 1D wave propagation
- Q&A regarding the final project

# 1/ Conceptual MATLAB example

- Diffusion of a step function between to conceptually distinct domains
- Communicate de appropriate inner boundary condition
- Repeat the exercise with the 1D wave equation
- Generalise it to n-processes
- Verify it against a single process implementation

## 2/ CUDA MPI 1D wave propagation

### CUDA-aware MPI - the important steps

- Initialise the context and get GPU ID from MPI

```
cudaSetDeviceFlags(cudaDeviceMapHost);  
const char* me_str      = getenv("OMPI_COMM_WORLD_RANK");  
const char* me_loc_str = getenv("OMPI_COMM_WORLD_LOCAL_RANK");  
me      = atoi(me_str);  
me_loc = atoi(me_loc_str);  
gpu_id = me_loc;
```



## 2/ CUDA MPI 1D wave propagation

### CUDA-aware MPI - the important steps

- MPI init functions; create a process world in cartesian coordinates, get neighbours

```
MPI_Init(&argc,&argv);
MPI_Comm_size(MPI_COMM_WORLD,&nprocs);
MPI_Dims_create(nprocs, NDIMS, dims);
MPI_Cart_create(MPI_COMM_WORLD, NDIMS, dims, periods, reorder, &topo_comm);
MPI_Comm_rank(topo_comm, &me);
MPI_Cart_coords(topo_comm, me, NDIMS, coords);
for (int i=0; i<NDIMS; i++){
    MPI_Cart_shift(topo_comm,i,1,&(neighbours_l[i]),&(neighbours_r[i]));
}
```

## 2/ CUDA MPI 1D wave propagation

### CUDA-aware MPI - the important steps

- MPI update sides
  - needs temporary buffers to store data to send (boundary values)
  - needs temporary buffers to store received data (boundary values)
  - needs to actually send and receive the data to be exchanged `update_sides()`

```
#define update_sides() \
if (neighbours_l[0] != MPI_PROC_NULL) write_to_mpi_sendbuffer_l0<<<grid,block>>>(V_send_l0_d,V_d);      cudaDeviceSynchronize(); \
if (neighbours_r[0] != MPI_PROC_NULL) write_to_mpi_sendbuffer_r0<<<grid,block>>>(V_send_r0_d,V_d,nx);  cudaDeviceSynchronize(); \
if (neighbours_l[0] != MPI_PROC_NULL){ MPI_Irecv((DAT*)V_recv_l0_d, 1, MPI_DAT, neighbours_l[0], tag, topo_comm, &(req[reqnr])); reqnr++; } \
if (neighbours_r[0] != MPI_PROC_NULL){ MPI_Irecv((DAT*)V_recv_r0_d, 1, MPI_DAT, neighbours_r[0], tag, topo_comm, &(req[reqnr])); reqnr++; } \
if (neighbours_l[0] != MPI_PROC_NULL){ MPI_Isend((DAT*)V_send_l0_d, 1, MPI_DAT, neighbours_l[0], tag, topo_comm, &(req[reqnr])); reqnr++; } \
if (neighbours_r[0] != MPI_PROC_NULL){ MPI_Isend((DAT*)V_send_r0_d, 1, MPI_DAT, neighbours_r[0], tag, topo_comm, &(req[reqnr])); reqnr++; } \
MPI_Waitall(reqnr,req,MPI_STATUSES_IGNORE); reqnr=0; for (int j=0; j<NREQS; j++){ req[j]=MPI_REQUEST_NULL; }; \
if (neighbours_l[0] != MPI_PROC_NULL) read_from_mpi_recvbuffer_l0<<<grid,block>>>(V_d, V_recv_l0_d);      cudaDeviceSynchronize(); \
if (neighbours_r[0] != MPI_PROC_NULL) read_from_mpi_recvbuffer_r0<<<grid,block>>>(V_d, V_recv_r0_d, nx); cudaDeviceSynchronize();
```

## 2/ CUDA MPI 1D wave propagation

### CUDA-aware MPI - further considerations

- Buffers can live in device memory only in case of CUDA-aware MPI
- Using `Isend` and `Irecv` for nonblocking message transmission > better performance
- Always start receiving first, then sending to avoid message losses and time loss

# Session 5 - CUDA MPI and Q&A

## Today's agenda

- Lecture: Multi-GPU implementation with CUDA MPI
- Programming: 1/ Conceptual MATLAB example  
2/ CUDA MPI 1D wave propagation
- Q&A regarding the final project

# Projects

Projects due date: August 1, 2019 - (midnight Pacific time)

- Objective: a/ 3D viscous Stokes flow: Buoyancy driven rising sphere setup  
 - or -  
 b/ 3D elastic wave propagation: initial gaussian pressure anomaly
- Results: 1/ Report performance (MTP), convergence with resolution increase  
 2a/ Viscous stokes: add convergence check to your code  
 2b/ Elastic (acoustic) wave: add P and S wave recording in a place of the domain
- Nice plots and fancy 3D graphics is a plus
- Hand in a report (max 5 pages) including: 1/ intro, 2/ motivation, 3/ mathematical model, 4/ numerical approach, 5/ results, 6/ discussion (personal thoughts on the topic) and conclusion.
- Hand in a zip file including the codes used to generate the results.

# Overall course conclusion

- Elastic waves and viscous Stokes flow
- Shared memory parallelisation: C CUDA and GPUs
- Stencil codes
- Convergence acceleration using second order scheme (damping)
- Distributed memory parallelisation: MPI + CUDA for multi-GPU configurations

# Suggested references

- Performance of stencil codes + MPI

[https://on-demand.gputechconf.com/gtc/2019/video/\\_/S9368/](https://on-demand.gputechconf.com/gtc/2019/video/_/S9368/)

- Iterative method for solving large 3D problems on GPUs

<http://www.nature.com/articles/s41598-018-29485-5>

<https://doi.org/10.1093/gji/ggz239>

<https://doi.org/10.1093/gji/ggy434>



● ● ● [wp.unil.ch/geocomputing/](http://wp.unil.ch/geocomputing/)

● ● ● [Iraess@stanford.edu](mailto:Iraess@stanford.edu)



